

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
c:\Python312\Lib\site-packages\seaborn\_statistics.py:32: UserWarning: A NumPy version >=1.22.4 and <2.3.0 is required for this version of SciPy (detected version 2.3.2)
from scipy.stats import gaussian_kde
```

```
In [2]: def read_csv_auto(path: str):
    for enc in ("utf-8", "utf-8-sig", "cp1252", "latin1"):
        try:
            df = pd.read_csv(path, encoding=enc)
            print(f"Loaded {path} with encoding={enc}")
            return df
        except UnicodeDecodeError:
            continue
    # Last resort: ignore undecodable chars
    df = pd.read_csv(path, encoding="latin1", encoding_errors="ignore")
    print(f"Loaded {path} with encoding=latin1 (errors=ignore)")
    return df

# Load datasets with robust encoding handling
stage_data = read_csv_auto('stage_data.csv')
tdf_stages = read_csv_auto('tdf_stages.csv')
tdf_winners = read_csv_auto('tdf_winners.csv')
```

```
Loaded stage_data.csv with encoding=utf-8
Loaded tdf_stages.csv with encoding=utf-8
Loaded tdf_winners.csv with encoding=utf-8
```

## Análisis Exploratorio de Datos (EDA)

En esta sección se realiza un EDA de los datasets cargados ( `stage_data` , `tdf_stages` , `tdf_winners` ).

```
In [3]: from typing import Optional, List

def quick_eda(df: pd.DataFrame, name: str,
              max_num_cols: int = 6,
              max_cat_cols: int = 6,
              top_n_categories: int = 10,
              corr_max_cols: int = 20,
              figsize=(12, 4)):
    """Genera un EDA compacto para el DataFrame dado y dibuja visuales básicos

    Args:
        df: DataFrame a analizar
        name: Nombre lógico del dataset
        max_num_cols: Máximo de columnas numéricas a graficar (histogramas)
        max_cat_cols: Máximo de columnas categóricas a graficar (barras)
```

```

        top_n_categories: Top N categorías a mostrar en barras
        corr_max_cols: Máximo de columnas numéricas en el mapa de correlación
        figsize: Tamaño base de las figuras
    """
    import warnings
    warnings.filterwarnings("ignore")

    print(f"\n===== {name} =====")
    print(f"Forma (filas, columnas): {df.shape}")
    print(f"\nTipos de datos:\n", df.dtypes)

    # Valores faltantes
    na_counts = df.isna().sum()
    na_pct = (na_counts / len(df) * 100).round(2)
    if na_counts.sum() > 0:
        na_table = pd.DataFrame({"na_count": na_counts, "na_pct": na_pct})
        print(f"\nValores faltantes (no-cero):\n", na_table[na_table.na_count > 0])
    else:
        print(f"\nValores faltantes: ninguno")

    # Duplicados
    dup_count = df.duplicated().sum()
    print(f"\nFilas duplicadas: {dup_count}")

    # Describe
    num_df = df.select_dtypes(include=np.number)
    if num_df.shape[1] > 0:
        desc = num_df.describe().T
        print(f"\nResumen numérico (describe):\n", desc)
        # Sesgo/curtosis
        sk = num_df.skew(numeric_only=True)
        ku = num_df.kurtosis(numeric_only=True)
        sk_ku = pd.concat([sk.rename("sesgo"), ku.rename("curtosis")], axis=1)
        print(f"\nSesgo/curtosis:\n", sk_ku)
    else:
        print(f"\nNo se detectaron columnas numéricas.")

    # Categóricas
    cat_cols = df.select_dtypes(include=["object", "category"]).columns.tolist()
    if cat_cols:
        print(f"\nColumnas categóricas y cardinalidad:")
        for c in cat_cols:
            nunique = df[c].nunique(dropna=True)
            print(f"  - {c}: {nunique} únicas")
    else:
        print(f"\nNo se detectaron columnas categóricas.")

    # Gráficos
    try:
        sns.set_theme(style="whitegrid")

        num_cols = num_df.columns.tolist()
        if num_cols:
            plot_cols = num_cols[:max_num_cols]
            n = len(plot_cols)
            if n:

```

```

fig, axes = plt.subplots(1, n, figsize=(figsize[0]*n/3, figs
if n == 1:
    axes = [axes]
for ax, col in zip(axes, plot_cols):
    sns.histplot(df[col].dropna(), kde=True, ax=ax)
    ax.set_title(f"Distribución: {col}")
plt.tight_layout()
plt.show()

if cat_cols:
    plot_cols = cat_cols[:max_cat_cols]
    n = len(plot_cols)
    if n:
        fig, axes = plt.subplots(1, n, figsize=(figsize[0]*n/3, figs
        if n == 1:
            axes = [axes]
        for ax, col in zip(axes, plot_cols):
            vc = df[col].value_counts(dropna=True).head(top_n_catego
            sns.barplot(x=vc.values, y=vc.index, ax=ax, orient="h")
            ax.set_title(f"Top {top_n_categories} {col}")
        plt.tight_layout()
        plt.show()

# Correlación
if num_cols:
    corr_cols = num_cols[:corr_max_cols]
    corr = df[corr_cols].corr(numeric_only=True)
    plt.figure(figsize=(min(1+0.5*len(corr_cols), 14), min(1+0.5*len
    sns.heatmap(corr, cmap="coolwarm", center=0, annot=False)
    plt.title("Matriz de correlación")
    plt.tight_layout()
    plt.show()

# Tendencia temporal si hay columna de año/fecha
def to_year_series(s):
    try:
        return pd.to_numeric(s, errors='coerce')
    except Exception:
        pass
    try:
        return pd.to_datetime(s, errors='coerce').dt.year
    except Exception:
        return pd.Series([np.nan]*len(s), index=s.index)

year_candidates = [c for c in df.columns if c.lower() in ("year", "a
if year_candidates and num_cols:
    ycol = year_candidates[0]
    yseries = to_year_series(df[ycol])
    ydf = pd.concat([yseries.rename("__year__"), num_df], axis=1)
    ydf = ydf.dropna(subset=["__year__"]) # Asegurar años válidos
    if not ydf.empty:
        agg = ydf.groupby("__year__")[num_cols].mean(numeric_only=Tr
        ax = agg.plot(figsize=(10,4), title=f"Promedios por año ({yc
        ax.set_xlabel("Año")
        plt.tight_layout()
        plt.show()

```

```
except Exception as e:  
    print("Se produjo un error en la sección de gráficos:", e)
```

```
In [4]: # Ejecutar EDA sobre cada dataset  
  
datasets = {  
    "stage_data": stage_data,  
    "tdf_stages": tdf_stages,  
    "tdf_winners": tdf_winners,  
}  
  
for name, df in datasets.items():  
    if isinstance(df, pd.DataFrame) and not df.empty:  
        quick_eda(df, name)  
    else:  
        print(f"Skipping {name}: not a valid non-empty DataFrame")
```

```
===== stage_data =====
Forma (filas, columnas): (255752, 11)
```

Tipos de datos:

```
edition          int64
year             int64
stage_results_id object
rank            object
time            object
rider           object
age             float64
team            float64
points          float64
elapsed         object
bib_number      object
dtype: object
```

Valores faltantes (no-cero):

	na_count	na_pct
team	255752	100.00
bib_number	254096	99.35
points	222746	87.09
time	5617	2.20
elapsed	5617	2.20
age	3326	1.30

Filas duplicadas: 0

Resumen numérico (describe):

	count	mean	std	min	25%	50%	75%
edition	255752.0	66.496950	26.817165	1.0	47.0	71.0	89.0
year	255752.0	1978.342930	28.939706	1903.0	1960.0	1984.0	2002.0
age	252426.0	27.547257	3.636306	13.0	25.0	27.0	30.0
team	0.0	NaN	NaN	NaN	NaN	NaN	NaN
points	33006.0	26.938526	27.035525	1.0	6.0	18.0	40.0

	max
edition	106.0
year	2019.0
age	49.0
team	NaN
points	100.0

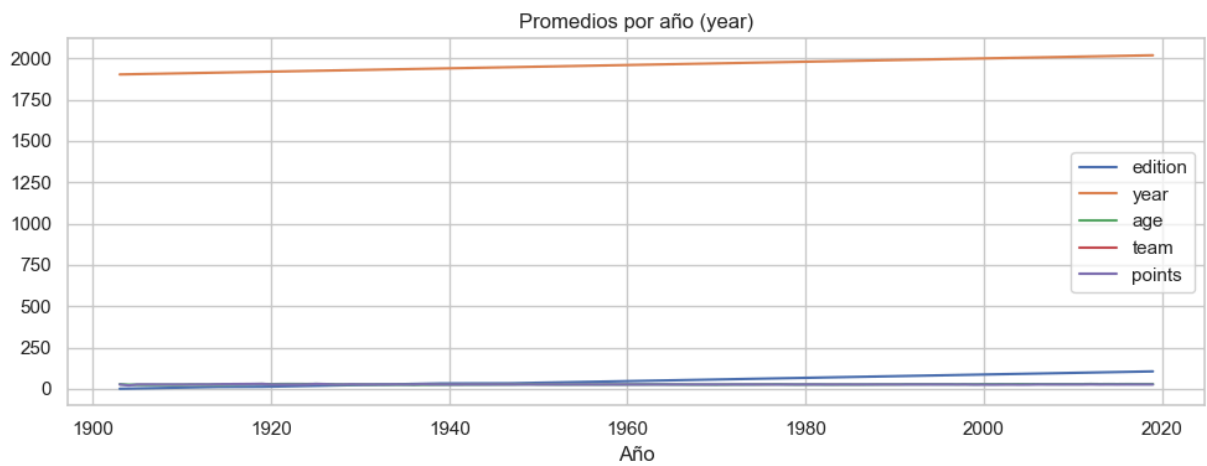

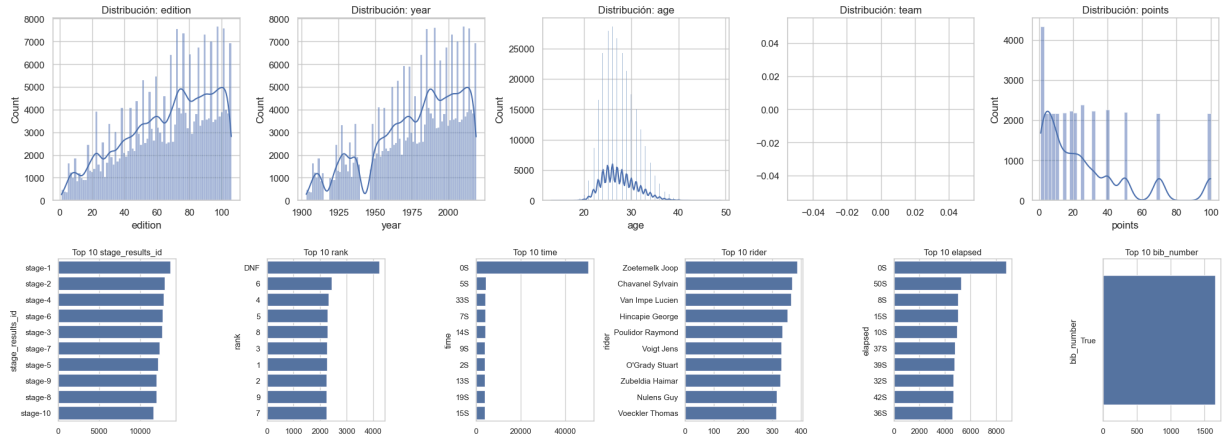
Sesgo/curtosis:

	sesgo	curtosis
edition	-0.461071	-0.760541
year	-0.646321	-0.433981
age	0.576350	0.178995
team	NaN	NaN
points	1.386451	1.165421

Columnas categóricas y cardinalidad:

- stage\_results\_id: 67 únicas
- rank: 219 únicas
- time: 60 únicas

- rider: 5162 únicas
- elapsed: 60 únicas
- bib\_number: 1 únicas



===== tdf\_stages =====

Forma (filas, columnas): (2236, 8)

Tipos de datos:

Stage	object
Date	object
Distance	float64
Origin	object
Destination	object
Type	object
Winner	object
Winner_Country	object
dtype:	object

Valores faltantes (no-cero):

	na_count	na_pct
Winner_Country	52	2.33

Filas duplicadas: 0

Resumen numérico (describe):



	count	mean	std	min	25%	50%	75%	max
Distance	2236.0	196.782994	90.176385	1.0	156.0	199.0	236.0	482.0

Sesgo/curtosis:


	sesgo	curtosis
Distance	0.16731	0.524551

Columnas categóricas y cardinalidad:

- Stage: 80 únicas
- Date: 2113 únicas
- Origin: 591 únicas
- Destination: 514 únicas
- Type: 18 únicas
- Winner: 878 únicas
- Winner\_Country: 40 únicas



Matriz de correlación





```
===== tdf_winners =====
Forma (filas, columnas): (106, 19)
```

Tipos de datos:

```
edition          int64
start_date       object
winner_name      object
winner_team      object
distance         float64
time_overall     float64
time_margin      float64
stage_wins       int64
stages_led       int64
height           float64
weight           float64
age              int64
born             object
died             object
full_name        object
nickname         object
birth_town       object
birth_country    object
nationality      object
dtype: object
```

Valores faltantes (no-cero):

	na_count	na_pct
full_name	60	56.60
died	50	47.17
height	40	37.74
weight	39	36.79
nickname	32	30.19
time_margin	8	7.55
time_overall	8	7.55

Filas duplicadas: 0

Resumen numérico (describe):

	count	mean	std	min	25%	\
edition	106.0	53.500000	30.743563	1.000000	27.250000	
distance	106.0	4212.064151	704.284160	2428.000000	3657.875000	
time_overall	98.0	125.754983	41.559391	82.086667	92.601597	
time_margin	98.0	0.267727	0.476194	0.002222	0.050833	
stage_wins	106.0	2.735849	1.842885	0.000000	1.000000	
stages_led	106.0	10.792453	5.307169	1.000000	6.250000	
height	66.0	1.778788	0.056989	1.610000	1.740000	
weight	67.0	69.253731	6.592795	52.000000	64.500000	
age	106.0	27.716981	3.354470	19.000000	26.000000	

	50%	75%	max
edition	53.500000	79.750000	106.000000
distance	4155.500000	4652.500000	5745.000000
time_overall	115.026806	142.678472	238.740278
time_margin	0.101667	0.249931	2.989167
stage_wins	2.000000	4.000000	8.000000
stages_led	12.000000	14.000000	22.000000

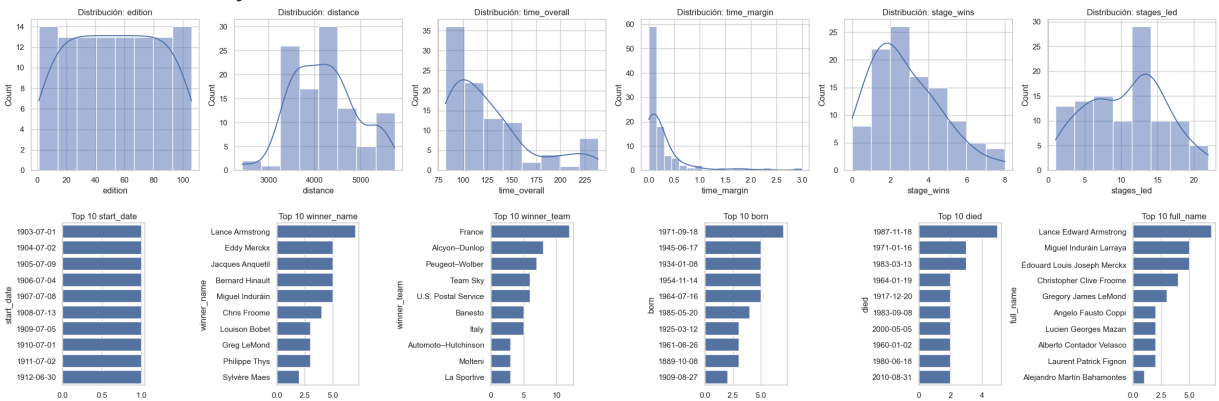
height	1.770000	1.820000	1.900000
weight	69.000000	74.000000	88.000000
age	28.000000	30.000000	36.000000


Sesgo/curtosis:

	sesgo	curtosis
edition	0.000000	-1.200000
distance	0.161691	-0.313684
time_overall	1.244370	0.713806
time_margin	3.578018	14.667470
stage_wins	0.723573	0.142827
stages_led	-0.011145	-0.822071
height	-0.402901	0.879929
weight	0.029010	0.225412
age	-0.029941	-0.420524

Columnas categóricas y cardinalidad:

- start\_date: 106 únicas
- winner\_name: 63 únicas
- winner\_team: 48 únicas
- born: 63 únicas
- died: 38 únicas
- full\_name: 23 únicas
- nickname: 37 únicas
- birth\_town: 58 únicas
- birth\_country: 15 únicas
- nationality: 14 únicas





## Conclusiones

### Panorama general

- Los tres datasets se cargaron con detección de codificación robusta y el EDA produjo resúmenes y visualizaciones útiles.
- Se observaron asimetrías y posibles outliers en variables numéricas.
- La correlación entre variables numéricas es moderada; no sugiere multicolinealidad severa en general.
- Donde hubo columna de año, hay tendencias temporales que justifican features por época o interacciones con tipo de etapa.

### Conclusiones por dataset

- stage\_data

- Tabla granular por etapa con métricas de desempeño; clave para entender diferencias por tipo de etapa.
- Útil para construir agregados por año/edición y relacionarlos con ganadores.
- tdf\_winners
  - Resumen a nivel ganador/edición.

This notebook was converted with [convert.ploomber.io](https://convert.ploomber.io)